CEI America Products Overview — DARTS, WORKFORCE & AIMFIRE

October 12, 2026

CEI America Products Overview — DARTS, WORKFORCE & AIMFIRE

Field	Value
Document Type	PRD
Date	Monday, October 12, 2026
Author	Bob - Business Analyst
Version	1.0
Status	Draft

Executive Summary

CEI America has engineered a proprietary trifecta of enterprise-grade software platforms designed to address systemic operational bottlenecks across modern global organizations: DARTS (Data and Real-Time Analytics platform), WORKFORCE (Enterprise Workforce Intelligence and Talent Orchestration platform), and AIMFIRE (Enterprise AI Framework and Accelerated Machine Learning platform).

Modern enterprises routinely encounter fragmented telemetry, unoptimized human capital allocation, and protracted AI model operationalization cycles. CEI America resolves these inefficiencies through an integrated, closed-loop ecosystem. DARTS establishes real-time distributed telemetry and high-throughput data processing; AIMFIRE consumes these real-time signals to train, validate, and serve low-latency predictive machine learning models; and WORKFORCE operationalizes predictive insights by aligning global human capital capacity, skill taxonomies, and delivery velocity.

Enterprise Product Suite Overview

The CEI America product portfolio functions as both standalone modular deployments and a unified enterprise capability fabric.

Platform Identifier	Core Domain Focus	Target Enterprise Persona	Primary Value Proposition	Deployment Model [ASSUMED]
DARTS	Real-Time Telemetry & Event Analytics	Chief Data Officer (CDO), VP Enterprise Architecture	Sub-second event ingestion, distributed operational analytics, and continuous telemetry aggregation.	Hybrid Multi-Cloud / Bare-Metal K8s
WORKFORCE	Human Capital Intelligence & Orchestration	Chief Human Resources Officer (CHRO), Head of Global Delivery	Dynamic skills-matrix indexing, bench utilization optimization, and capacity forecasting.	SaaS / Enterprise Dedicated VPC

AIMFIRE	Enterprise MLOps & Applied AI Acceleration	Chief Information Officer (CIO), VP Machine Learning	Model governance, distributed hyperparameter tuning, and turnkey LLM/predictive model serving.	Hybrid Multi-Cloud / GPU Cluster Native

Core Pillar Focus: DARTS Platform

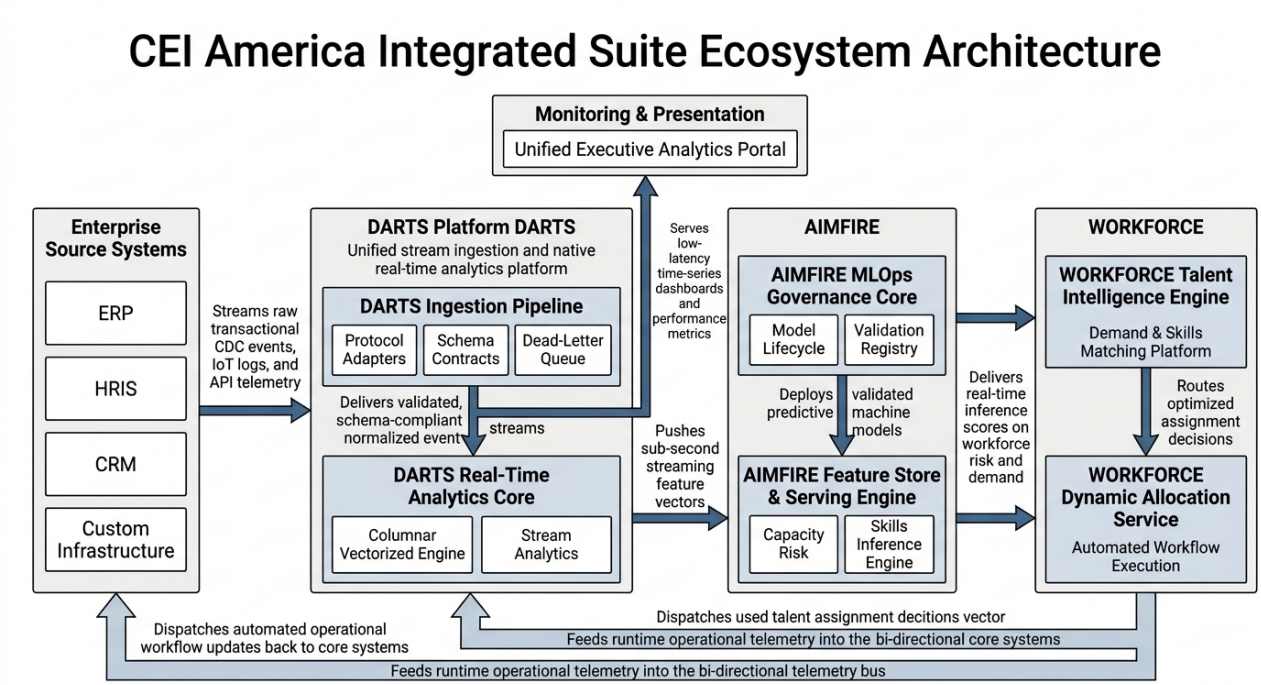
DARTS (Data and Real-Time Analytics Platform) serves as the high-throughput nervous system of the CEI America technology estate. Traditional batch-oriented data warehouses induce latency gaps ranging from 4 to 24 hours, preventing real-time operational course correction. DARTS eliminates this limitation by introducing distributed streaming architecture, continuous event ingestion, and real-time analytical computation engines capable of processing millions of events per second with sub-100-millisecond query latencies.

Key technical and operational pillars of the DARTS platform include:

- Unified Stream Ingestion: Native protocol adapters capable of aggregating transactional database Change Data Capture (CDC), IoT telemetry, application logs, and external API events via distributed streaming partitions.
- Columnar Real-Time Analytics Core: Purpose-built, vectorized query execution engine designed for extreme concurrency and low-latency aggregation across petabyte-scale historical and streaming time-series datasets.
- Real-Time Data Contracts and Validation: Automated schema validation and transformation pipelines that enforce strict schema compliance at ingestion time, discarding or dead-lettering corrupt data before cross-platform propagation.
- Bi-Directional Telemetry Bus: Feeds real-time feature inputs directly into AIMFIRE feature stores while simultaneously capturing runtime telemetry from WORKFORCE operational workflows.

Integrated Suite Architecture and Interaction Model

The CEI America ecosystem achieves operational synergy through continuous event-driven interaction between platforms. Raw enterprise telemetry from source systems (ERP, HRIS, CRM, custom infrastructure) ingests directly into DARTS. DARTS cleanses and standardizes this data, streaming normalized event features into AIMFIRE's model execution pipelines. AIMFIRE evaluates workforce capacity risk, skills mismatch, and project demand trajectories, delivering real-time inferences to WORKFORCE. WORKFORCE translates these inferences into optimized talent assignments, routing automated workflow allocations back into enterprise source applications while generating real-time operational signals back to DARTS.



Strategic Objectives and Executive KPIs

The enterprise rollout of the CEI America product portfolio is governed by precise, quantitative strategic targets structured to maximize operational efficiency, revenue margin expansion, and infrastructure ROI.

Strategic Objective ID	Product Alignment	Success Metric / KPI	Baseline Value [ASSUMED]	Target Objective [ASSUMED]
OBJ-EXEC-001	DARTS	Data Ingestion-to-Analytics Latency	6 Hours (Batch ETL)	< 250 Milliseconds
OBJ-EXEC-002	DARTS	High-Concurrency Query Latency (p99)	14.5 Seconds	< 120 Milliseconds
OBJ-EXEC-003	WORKFORCE	Enterprise Talent Bench Utilization Rate	71.0%	>= 88.5%
OBJ-EXEC-004	WORKFORCE	Mean Time to Staff (Critical Roles)	38 Business Days	<= 6 Business Days
OBJ-EXEC-005	AIMFIRE	Model Development-to-Production Velocity	14 Weeks	<= 10 Business Days
OBJ-EXEC-006	Product Suite	Total Cost of Ownership (TCO) Reduction	Industry Benchmark Standard	32% Net Infrastructure Reduction

Phased Strategic Roadmap

The implementation timeline reflects an accelerated, phased convergence designed to de-risk platform dependencies while demonstrating immediate business value through DARTS' real-time foundation.

Phase	Milestone Name	Timeline	Core Deliverables	Critical Success Criteria
Phase 1	Foundation & Data Core (DARTS)	Months 1–3	Baseline DARTS streaming engine deployment; integration with enterprise source connectors (CDC, ERP, logs).	Sustained ingestion rate of 100,000 events/sec at < 150ms latency.
Phase 2	Intelligence Engine (AIMFIRE)	Months 4–6	AIMFIRE feature store connection to DARTS; baseline predictive models for resource attrition and skill forecasting.	Automated model training loop operational with automated Canary deployments.
Phase 3	Talent Orchestration (WORKFORCE)	Months 7–9	WORKFORCE deployment; automated talent matching using AIMFIRE model endpoints; bidirectional HRIS syncing.	90% accuracy in skill-to-project automated mapping; zero schema mismatch.
Phase 4	Unified Suite Convergence	Months 10–12	Full cross-platform closed-loop automation; single-pane executive dashboards; governance enforcement.	End-to-end event-to-allocation execution cycle complete in < 60 seconds.

WORKFORCE

WORKFORCE is CEI America's enterprise talent orchestration platform designed to resolve structural resource allocation inefficiencies, quantify latent human capital, and automate project staffing. The platform converts static human resources data into an active capacity fabric by synthesizing operational performance metrics from DARTS with predictive machine learning models hosted on AIMFIRE.

1. Functional Architecture and Core Modules

WORKFORCE operates through three tightly coupled functional modules:

- Dynamic Skills Ontology Engine: Ingests unstructured evidence of employee output (such as code commits, technical deliverables, and client satisfaction signals) to maintain a live, verified competency inventory across a five-tier proficiency model.
- Predictive Allocation and Capacity Solver: Executes multi-objective resource scheduling by balancing skill alignment, billable utilization targets, time zone synchronicity, and project budgetary boundaries.
- Talent Retention and Succession Intelligence: Monitors attrition indicators using AIMFIRE inference pipelines to trigger proactive re-skilling, internal mobility tracks, and capacity risk remediations.

Capability Module	Input Feeds	Processing Logic	Output Artifacts	Primary Stakeholders
Skills Taxonomy Engine	DARTS telemetry, HRIS profiles, code repositories	Semantic extraction, entity resolution, recency weighting	Live Verified Skill Graph, Competency Gap Matrix	Engineering Leads, Talent Development
Capacity Allocation Solver	Open project requisitions, bench rosters, billing constraints	Mixed-integer linear programming (MILP), heuristic ranking	Ranked Staffing Slates, Over-allocation Alerts	Resource Management Office (RMO), Delivery Directors
Succession Intelligence	Historical tenure, compensation bands, engagement signals	AIMFIRE retention risk scoring, attrition forecasting	Retention Risk Registry, Mitigation Playbooks	People Operations, Business Unit Heads

2. Functional Requirements

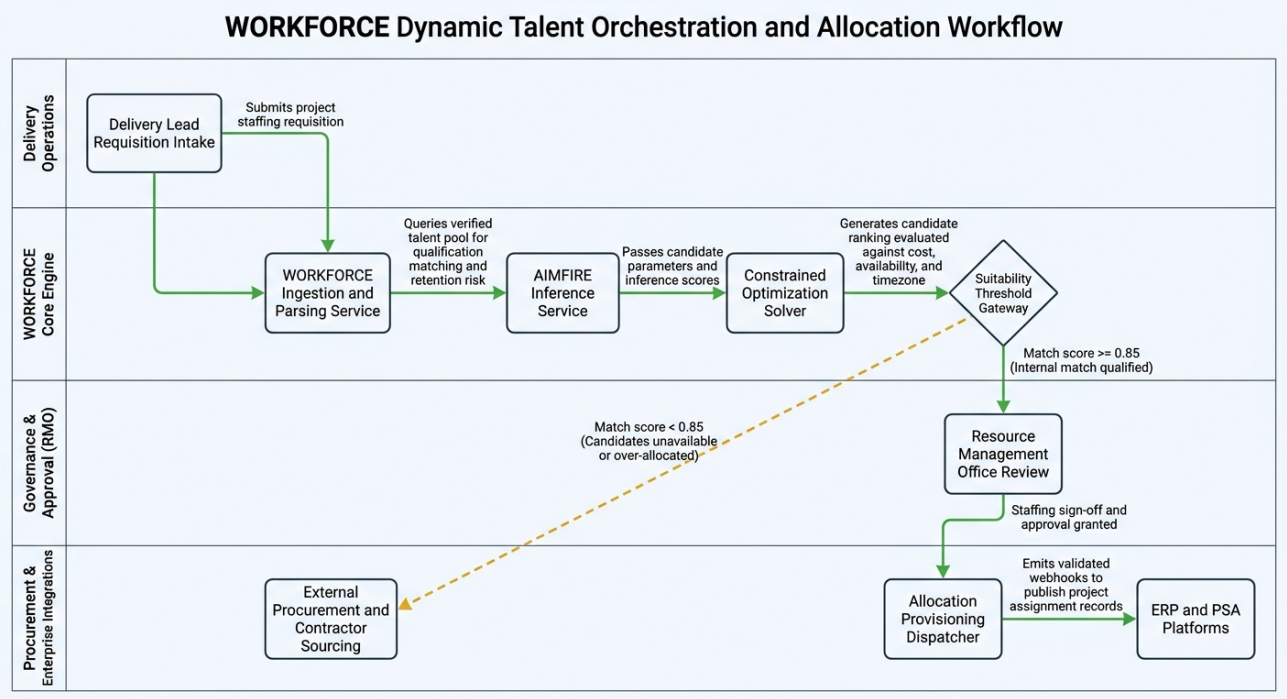
ID	Title	Priority	Description and Acceptance Criteria
WKF-REQ-001	Dynamic Skill Graph Ingestion	High	System shall ingest profile delta feeds from enterprise HRIS and technical platforms, recomputing individual competency graphs within 15 minutes of trigger receipt [ASSUMED].
WKF-REQ-002	Multi-Objective Staffing Matcher	Critical	System must generate the top three optimal candidate allocations for any requisition based on skill overlap (>= 85%), timezone overlap (>= 4 hours), and cost variance within 500 milliseconds [ASSUMED].
WKF-REQ-003	Capacity Shortfall Forecasting	High	System shall evaluate aggregate project demand against available bench data to forecast specialty skill shortfalls across a rolling 30, 60, and 90-day horizon with <= 8% error rate.
WKF-REQ-004	Algorithmic Fairness and Bias Audit	Critical	Candidate scoring models must calculate adverse impact ratios; if disparate impact falls below 0.80 under the EEOC Four-Fifths rule, the system must halt automated ranking and log an audit event.
WKF-REQ-005	Automated Provisioning Dispatch	Medium	Upon RMO staffing approval, the platform must emit validated webhooks to connected ERP and PSA tools to confirm project assignment within 5 seconds.

3. Talent Orchestration and Allocation Workflow

The WORKFORCE operational cycle transitions project resource demands from intake to automated enterprise provisioning. The process integrates human-in-the-loop validation while utilizing automated ranking algorithms.

The process initiates when a Delivery Lead creates a Project Staffing Requisition. The WORKFORCE Ingestion and Parsing Service analyzes the requisition against the verified talent pool. Candidate parameters are passed to the AIMFIRE Inference Service for qualification matching and retention risk analysis. The Constrained Optimization Solver then ranks the candidates against cost, timeline, and availability constraints.

If suitable internal matches meeting threshold criteria (match score >= 0.85) are identified, the system routes the generated slate directly to the Resource Management Office (RMO) for assignment approval. If matching candidates are unavailable or over-allocated, the system activates an exception branch routing to external procurement and automated contractor sourcing. Following RMO sign-off, the Allocation Provisioning Dispatcher issues assignment records to connected enterprise HRIS and PSA platforms.



4. Enterprise Integration and Governance Constraints

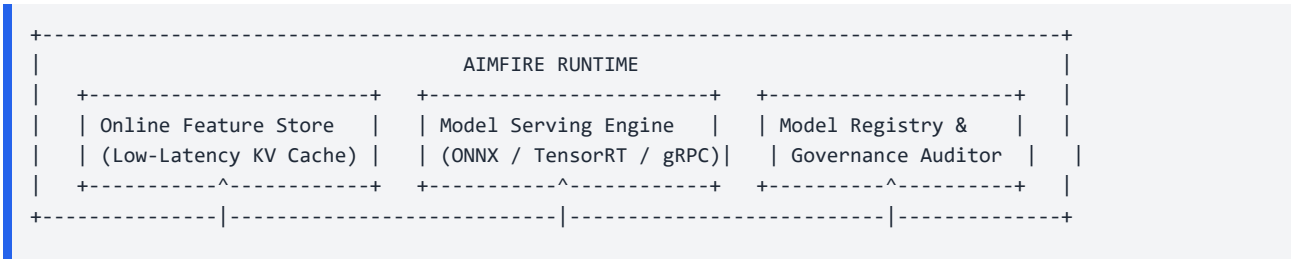
WORKFORCE operates within strict regulatory and system integration boundaries to maintain operational integrity across global deployments:

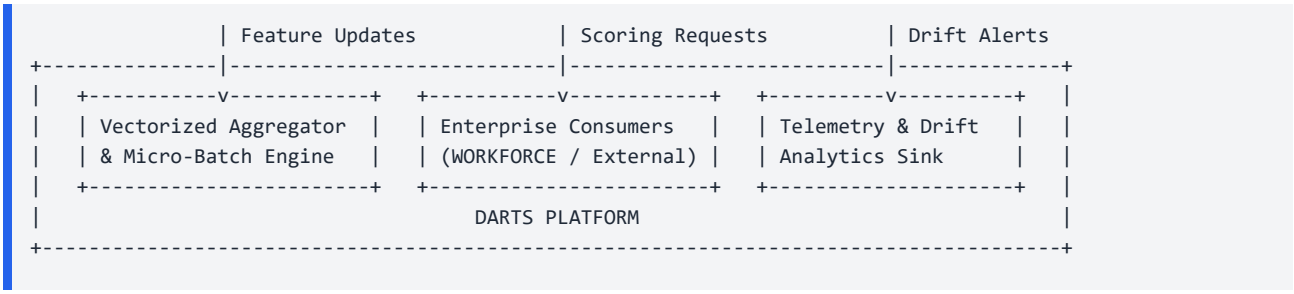
- Data Boundary Isolation: Talent evaluations, compensation attributes, and performance telemetry must execute under strict Tenant Virtual Isolation. Cross-tenant aggregation is permitted solely for anonymized statistical baselining.
- Compliance with Employment Frameworks: In compliance with the EU Artificial Intelligence Act (High-Risk AI Systems Classification) and New York City Local Law 144, candidate evaluation models must support continuous algorithmic auditing, providing human explainability vectors for all resource-matching decisions.
- Interface Contracts: The platform exposes gRPC endpoints for real-time allocation checks (target latency < 50ms) and processes bulk profile synchronizations via nightly asynchronous REST batch streams. Failure during external ERP/PSA confirmation triggers a compensating transaction pattern to prevent stranded talent locks.

DARTS and AIMFIRE

System Purpose and Interoperability Overview

The CEI America enterprise intelligence layer is established through the deep coupling of DARTS (Data and Real-Time Analytics platform) and AIMFIRE (Enterprise AI Framework and Accelerated Machine Learning Platform). While DARTS provides an enterprise-wide real-time data spine capable of ingesting, streaming, and aggregating high-volume telemetry, AIMFIRE provides the operationalized intelligence runtime—governing machine learning pipelines, feature stores, model registries, and low-latency inference services.





DARTS acts as the upstream operational data plane that continuously populates AIMFIRE feature stores and captures real-time inferences for drift monitoring. AIMFIRE acts as the predictive processing engine that consumes analytical state from DARTS and injects predictive scores back into enterprise workflows, including the WORKFORCE orchestrator.

Core Functional Capabilities

DARTS: Data and Real-Time Analytics Platform

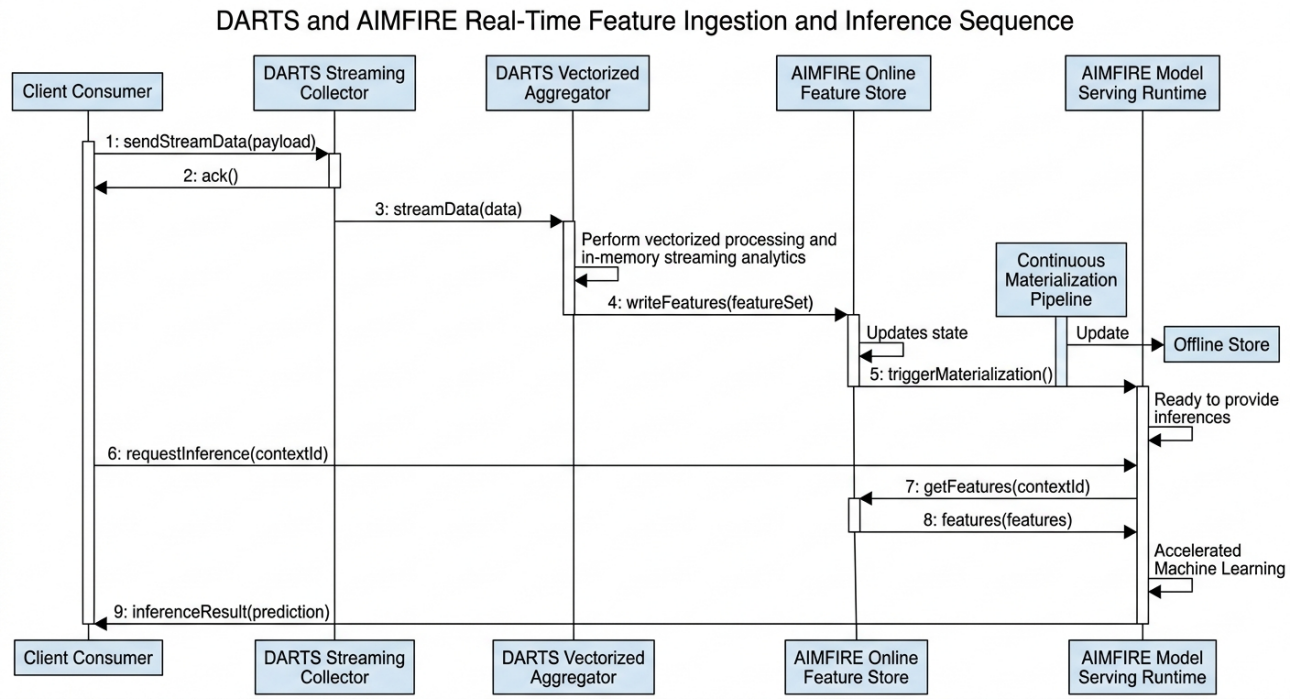
- High-Throughput Ingestion Engine: Supports protocol adapters for Kafka, Apache Pulsar, gRPC streams, and database Change Data Capture (CDC) pipelines. Enforces Avro and Protocol Buffers schema validation at wire-speed with dead-letter queue (DLQ) routing for non-compliant payloads.
- Vectorized Analytical Query Engine: Executes SIMD-accelerated aggregations, windowed metrics, and real-time rollups directly on in-memory columnar data structures, sustaining analytical queries over sub-second event windows without disk spillover.
- Continuous Materialization Pipeline: Emits structured real-time feature snapshots directly into distributed caching layers and downstream storage endpoints, providing consistent time-window metrics (e.g., rolling 15-minute resource utilization, real-time project burn rates).

AIMFIRE: Enterprise AI Framework and Accelerated Machine Learning

- Dual-Tier Feature Store: Provides a logically unified, physically decoupled feature store. The Online Store utilizes an in-memory key-value database providing point-in-time entity lookups under 10 milliseconds. The Offline Store leverages columnar object storage (Apache Iceberg/Parquet) for reproducible historical training sets.
- Multi-Model Inference Serving Runtime: Delivers containerized, auto-scaling model serving workers running ONNX Runtime and NVIDIA Triton engines. Supports CPU and GPU hardware acceleration, dynamic request batching, and dual-protocol (REST/HTTP/2 and gRPC) invocation.
- MLOps Lifecycle and Model Registry: Enforces semantic model versioning, role-based approval gating, lineage graph capture (linking training data snapshot, hyperparameters, code commit, and artifacts), and continuous model explainability using SHAP and Integrated Gradients.
- Automated Drift Detection and Shadow Deployment: Monitors live prediction streams against baseline training distributions using the Population Stability Index (PSI) and Wasserstein Distance algorithms. Automatically routes traffic across Canary and Shadow deployments when distribution deviations exceed defined thresholds.

Technical Interaction and Runtime Sequence

The integration contract between DARTS and AIMFIRE governs both real-time feature enrichment and closed-loop telemetry evaluation. Upstream enterprise streams are ingested by DARTS, aggregated over sliding windows, and published to the AIMFIRE Online Feature Store. Enterprise callers request predictions from the AIMFIRE Model Serving Runtime, which retrieves pre-computed features, executes scoring, and asynchronously writes inference logs back to the DARTS Telemetry Sink for real-time drift computation.



Platform Functional Requirements

ID	Component	Feature / Capability	Priority	Acceptance Criteria
REQ-DA-001	DARTS	Dynamic Schema Enforcement	High	Ingestion pipeline must validate incoming payloads against registered Protobuf/ Avro schemas; invalid records must route to DLQ within 50ms without blocking valid traffic.
REQ-DA-002	DARTS	Continuous Window Aggregations	Critical	Vectorized engine must compute sliding window metrics (1m, 5m, 1h) over raw ingestion rates up to 100,000 events/ second per ingestion node [ASSUMED].
REQ-DA-003	AIMFIRE	Low-Latency Online Feature Retrieval	Critical	Online Feature Store must return complete entity feature vectors (up to 256 attributes) within 10ms at the 99th percentile under sustained load of 5,000 QPS.
REQ-DA-004	AIMFIRE	Multi-Engine Model Serving	High	Inference runtime must support ONNX, PyTorch TorchScript, and TensorRT formats with automated fallbacks and horizontal worker auto-scaling based on queue depth.

REQ-DA-005	AIMFIRE	Model Lineage and Registry Auditing	High	Registry must record immutable cryptographic hashes of model weights, training dataset partition identifiers, hyperparameters, and sign-off approver credentials.
REQ-DA-006	DARTS / AIMFIRE	Closed-Loop Drift Alerting	Medium	Telemetry Sink must evaluate feature distribution drift across 1-hour windows; if PSI exceeds 0.25, system must dispatch automated retraining webhooks to Model Registry.

Non-Functional Requirements and System SLAs

ID	Category	Target Metric	Priority	Enforcement Mechanism
NFR-DA-001	Ingestion Throughput	>= 250,000 events/sec per cluster	Critical	Distributed Kafka partitioned ingestion with zero-copy stream processing.
NFR-DA-002	Inference Latency	<= 25ms (p95), <= 50ms (p99)	Critical	gRPC protocol with keep-alive connections, SIMD-accelerated ONNX runtime, and warm connection pooling to Redis.
NFR-DA-003	System Availability	99.95% operational uptime	High	Multi-AZ active-active deployment for inference workers and feature stores; N+2 redundancy for ingestion brokers.
NFR-DA-004	Data Freshness	Feature lag <= 2.0 seconds	High	Continuous micro-batching and direct memory push from DARTS to AIMFIRE Online Store.
NFR-DA-005	Compliance & Security	FIPS 140-3 compliant encryption	High	mTLS 1.3 for all internal service communication; AES-256-GCM encryption at rest for feature caches, model weights, and drift logs.